

# Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>  
(<https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

## Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

## [1]. Reading Data

### [1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

In [1]:

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

from sklearn.model_selection import train_test_split

from scipy.sparse import csr_matrix, issparse
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
from sklearn.metrics import classification_report
import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer
from sklearn.neighbors import KNeighborsClassifier
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
```

In [2]:

```
# using SQLite Table to read data.
con = sqlite3.connect('../input/database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data po
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIM
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a nega
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)
```

Number of data points in our data (30000, 10)

Out[2]:

		Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenon
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian		1	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa		0	
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"		1	

In [3]:

```
display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

In [4]:

```
print(display.shape)
display.head()
```

(80668, 7)

Out[4]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
0	#oc-R115TNMSPFT9I7	B005ZBZLT4	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJB9	B005HG9ESG	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B005ZBZLT4	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ESG	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBEV0	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

In [5]:

```
display[display['UserId']=='AZY10LLTJ71NX']
```

Out[5]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
80638	AZY10LLTJ71NX	B001ATMQK2	undertheshrine "undertheshrine"	1296691200	5	I bought this 6 pack because for the price tha...	5

In [6]:

```
display['COUNT(*)'].sum()
```

Out[6]:

393063

## [2] Exploratory Data Analysis

### [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

In [7]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

Out[7]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDen
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator,

HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

In [8]:

```
#Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=
```

In [9]:

```
#Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, ke
final.shape
```

Out[9]:

(28072, 10)

In [10]:

```
#Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[10]:

93.57333333333332

**Observation:-** It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

In [11]:

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

Out[11]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	

In [12]:

```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

In [13]:

```
#Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()

(28072, 10)
```

Out[13]:

```
1    23606
0     4466
Name: Score, dtype: int64
```

## [3] Preprocessing

### [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

In [14]:

```
# printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

Our dogs just love them. I saw them in a pet store and a tag was attached regarding them being made in China and it satisfied me that they were safe.

=====

When I ordered these, I thought they were a bit pricey, but I decided to give them a try anyway. I'm glad I did! My dogs absolutely love these dried liver treats. And, since my dogs are all small, I can cut the treats in half and still have large enough pieces to satisfy them. They're great for training; I'll definitely order them again, and would recommend them to anyone.

=====

This was my favorite stevia product and I had it on subscribe and save until I queried customer service about NuNaturals GMO use. Yes, NuNaturals uses GMO products. SO, I've canceled my subscribe and save order and am now using [Stevita Stevia Clear Liquid Extract, 3.3-Ounce Container \(Pack of 3\)](http://www.amazon.com/gp/product/B001ELL3U0).

=====

TOTALLY ORGASMIC. these chips are the best spicy chip i have ever tasted. signed up for the subscribe and save option. the case contained 15(FIFTEEN, FULL SIZED BAGS) OF CHIPS. the price per unit equals \$1.73 per package. that is not even the cost of plain chips. if you add the free shipping and the fast delivery, this deal is a steal. so run like you stole something over to your computer and order the SPICY THAI CHIPS. p.s. even if you paid the going price of \$30.00, you are still ahead of the curve. ENJOY

=====



In [15]:

```
# remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

Our dogs just love them. I saw them in a pet store and a tag was attached regarding them being made in China and it satisfied me that they were safe.

In [16]:

```
# https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-a
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)
```

Our dogs just love them. I saw them in a pet store and a tag was attached regarding them being made in China and it satisfied me that they were safe.

=====

When I ordered these, I thought they were a bit pricey, but I decided to give them a try anyway. I'm glad I did! My dogs absolutely love these dried liver treats. And, since my dogs are all small, I can cut the treats in half and still have large enough pieces to satisfy them. They're great for training; I'll definitely order them again, and would recommend them to anyone.

=====

This was my favorite stevia product and I had it on subscribe and save until I queried customer service about NuNaturals GMO use. Yes, NuNaturals uses GMO products. SO, I've canceled my subscribe and save order and am now using Stevita Stevia Clear Liquid Extract, 3.3-Ounce Container (Pack of 3).

=====

TOTALLY ORGASMIC. these chips are the best spicy chip i have ever tasted. signed up for the subscribe and save option. the case contained 15(FIFTEEN, FULL SIZED BAGS) OF CHIPS. the price per unit equals \$1.73 per package. that is not even the cost of plain chips. if you add the free shipping and the fast delivery, this deal is a steal. so run like you stole something over to your computer and order the SPICY THAI CHIPS. p.s. even if you paid the going price of \$30.00, you are still ahead of the curve. ENJOY

In [17]:

```
# https://stackoverflow.com/a/47091490/4084039
```

```
import re
```

```
def decontracted(phrase):
```

```
    # specific
```

```
    phrase = re.sub(r"won't", "will not", phrase)
```

```
    phrase = re.sub(r"can't", "can not", phrase)
```

```
    # general
```

```
    phrase = re.sub(r"n't", " not", phrase)
```

```
    phrase = re.sub(r"'re", " are", phrase)
```

```
    phrase = re.sub(r"'s", " is", phrase)
```

```
    phrase = re.sub(r"'d", " would", phrase)
```

```
    phrase = re.sub(r"'ll", " will", phrase)
```

```
    phrase = re.sub(r"'t", " not", phrase)
```

```
    phrase = re.sub(r"'ve", " have", phrase)
```

```
    phrase = re.sub(r"'m", " am", phrase)
```

```
    return phrase
```

In [18]:

```
sent_1500 = decontracted(sent_1500)
```

```
print(sent_1500)
```

```
print("="*50)
```

This was my favorite stevia product and I had it on subscribe and save until I queried customer service about NuNaturals GMO use. Yes, NuNaturals uses GMO products. SO, I have canceled my subscribe and save order and am now using <http://www.amazon.com/gp/product/B001ELL3U0> Stevita Stevia Clear Liquid Extract, 3.3-Ounce Container (Pack of 3).

=====

In [19]:

```
#remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
```

```
sent_0 = re.sub(r"\S*\d\S*", "", sent_0).strip()
```

```
print(sent_0)
```

Our dogs just love them. I saw them in a pet store and a tag was attached regarding them being made in China and it satisfied me that they were safe.

In [20]:

```
#remove spacial character: https://stackoverflow.com/a/5843547/4084039
```

```
sent_1500 = re.sub(r'[^A-Za-z0-9]+', ' ', sent_1500)
```

```
print(sent_1500)
```

This was my favorite stevia product and I had it on subscribe and save until I queried customer service about NuNaturals GMO use Yes NuNaturals uses GMO products SO I have canceled my subscribe and save order and am now using a href http www amazon com gp product B001ELL3U0 Stevita Stevia Clear Liquid Extract 3 3 Ounce Container Pack of 3 a

In [21]:

```
# https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourse',
               "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'hi',
               'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
               'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that',
               'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has',
               'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because',
               'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'thro',
               'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off',
               'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all',
               'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 't',
               's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've",
               've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
               "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mustn't',
               'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'won', "won't",
               'wouldn', "wouldn't"])
```

In [22]:

```
# Combining all the above stundents
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

100%|██████████| 28072/28072 [00:11<00:00, 2392.03it/s]

In [23]:

```
preprocessed_reviews[1500]
```

Out[23]:

```
'favorite stevia product subscribe save queried customer service nunat
urals gmo use yes nunaturals uses gmo products canceled subscribe save
order using'
```

## [3.2] Preprocessing Review Summary

In [24]:

```
## Similarly you can do preprocessing for review summary also.
```

## [4] Featurization

### [4.1] BAG OF WORDS

```
#BoW count_vect = CountVectorizer() #in scikit-learn count_vect.fit(preprocessed_reviews) print("some feature names ", count_vect.get_feature_names()[:10]) print('='*50)

final_counts = count_vect.transform(preprocessed_reviews) print("the type of count vectorizer ",type(final_counts)) print("the shape of out text BOW vectorizer ",final_counts.get_shape()) print("the number of unique words ", final_counts.get_shape()[1])
```

### [4.2] Bi-Grams and n-Grams.

4g#bi-gram, tri-gram and n-gram

#removing stop words like "not" should be avoided before building n-grams

**count\_vect = CountVectorizer(ngram\_range=(1,2))**

**please do read the CountVectorizer documentation**

[http://scikit-learn.org/stable/modules/generated/sklearn.feature\\_extraction.text.CountVectorizer.html](http://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html)  
([http://scikit-learn.org/stable/modules/generated/sklearn.feature\\_extraction.text.CountVectorizer.html](http://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html))

**you can choose these numebrs min\_df=10,  
max\_features=5000, of your choice**

```
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000) final_bigram_counts = count_vect.fit_transform(preprocessed_reviews) print("the type of count vectorizer ",type(final_bigram_counts)) print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape()) print("the number of unique words including both unigrams and bigrams ", final_bigram_counts.get_shape()[1])
```

Type *Markdown* and LaTeX:  $\alpha^2$

### [4.3] TF-IDF

```
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10) tf_idf_vect.fit(preprocessed_reviews) print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names()[0:10]) print('='*50)
```

```
final_tf_idf = tf_idf_vect.transform(preprocessed_reviews) print("the type of count vectorizer ",type(final_tf_idf)) print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape()) print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[1])
```

## [4.4] Word2Vec

### Train your own Word2Vec model using your own text corpus

```
i=0 list_of_sentence=[] for sentence in preprocessed_reviews: list_of_sentence.append(sentence.split())
```

### Using Google News Word2Vectors

in this project we are using a pretrained model by google

its 3.3G file, once you load this into your memory

it occupies ~9Gb, so please do this step only if you have >12G of ram

we will provide a pickle file wich contains a dict ,

and it contains all our courpus words as keys and model[word] as values

To use this code-snippet, download "GoogleNews-vectors-negative300.bin"

from

<https://drive.google.com/file/d/0B7XkCwpl5KDYNINUTTI9/view>  
(<https://drive.google.com/file/d/0B7XkCwpl5KDYNINUTTI9/view>)

it's 1.9GB in size.

<http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY> (<http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY>)

you can comment this whole cell

or change these variable according to your need

```
is_your_ram_gt_16g=False want_to_use_google_w2v = False want_to_train_w2v = True
```

```
if want_to_train_w2v:
```

```
    # min_count = 5 considers only words that occurred at least 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))
```

```
elif want_to_use_google_w2v and is_your_ram_gt_16g: if os.path.isfile('GoogleNews-vectors-negative300.bin'):
w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin', binary=True)
print(w2v_model.wv.most_similar('great')) print(w2v_model.wv.most_similar('worst')) else: print("you don't have
gogole's word2vec file, keep want_to_train_w2v = True, to train your own w2v ")
```

```
w2v_words = list(w2v_model.wv.vocab) print("number of words that occurred minimum 5 times
",len(w2v_words)) print("sample words ", w2v_words[0:50])
```

## **[4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V**

### **[4.4.1.1] Avg W2v**

## **average Word2Vec**

## **compute average word2vec for each review.**

```
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list for sent in tqdm(list_of_sentence):
# for each review/sentence sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to
change this to 300 if you use google's w2v cnt_words = 0; # num of words with a valid vector in the
sentence/review for word in sent: # for each word in a review/sentence if word in w2v_words: vec =
w2v_model.wv[word] sent_vec += vec cnt_words += 1 if cnt_words != 0: sent_vec /= cnt_words
sent_vectors.append(sent_vec) print(len(sent_vectors)) print(len(sent_vectors[0]))
```

### **[4.4.1.2] TFIDF weighted W2v**

```
S = ["abc def pqr", "def def def abc", "pqr pqr def"]
```

```
model = TfidfVectorizer() tf_idf_matrix = model.fit_transform(preprocessed_reviews)
```

**we are converting a dictionary with word as a key, and the idf as a value**

```
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

# TF-IDF weighted Word2Vec

```
tfidf_feat = model.get_feature_names() # tfidf words/col-names
```

**final\_tf\_idf is the sparse matrix with row= sentence, col=word and cell\_val = tfidf**

```
tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list row=0; for sent in
tqdm(list_of_sentence): # for each review/sentence sent_vec = np.zeros(50) # as word vectors are of zero
length weight_sum =0; # num of words with a valid vector in the sentence/review for word in sent: # for each
word in a review/sentence if word in w2v_words and word in tfidf_feat: vec = w2v_model.wv[word]
```

**tf\_idf = tf\_idf\_matrix[row, tfidf\_feat.index(word)]**

```
# to reduce the computation we are
# dictionary[word] = idf value of word in whole corpus
# sent.count(word) = tf value of word in this review
tf_idf = dictionary[word]*(sent.count(word)/len(sent))
sent_vec += (vec * tf_idf)
weight_sum += tf_idf
if weight_sum != 0:
    sent_vec /= weight_sum
tfidf_sent_vectors.append(sent_vec)
row += 1
```

## [5] Assignment 3: KNN

### 1. Apply Knn(brute force version) on these feature sets

- **SET 1:** Review text, preprocessed one converted into vectors using (BOW)
- **SET 2:** Review text, preprocessed one converted into vectors using (TFIDF)
- **SET 3:** Review text, preprocessed one converted into vectors using (AVG W2v)
- **SET 4:** Review text, preprocessed one converted into vectors using (TFIDF W2v)

### 2. Apply Knn(kd tree version) on these feature sets

**NOTE:** sklearn implementation of kd-tree accepts only dense matrices, you need to convert the sparse matrices of CountVectorizer/TfidfVectorizer into dense matrices. You can convert sparse matrices to dense using .toarray() attribute. For more information please visit this [link \(https://docs.scipy.org/doc/scipy-0.18.1/reference/generated/scipy.sparse.csr\\_matrix.toarray.html\)](https://docs.scipy.org/doc/scipy-0.18.1/reference/generated/scipy.sparse.csr_matrix.toarray.html).

- **SET 5:** Review text, preprocessed one converted into vectors using (BOW) but with restriction on maximum features generated.

```
count_vect = CountVectorizer(min_df=10, max_features=50
0)
count_vect.fit(preprocessed_reviews)
```



- **SET 6:** Review text, preprocessed one converted into vectors using (TFIDF) but with restriction on maximum features generated.

```
tf_idf_vect = TfidfVectorizer(min_df=10, max_feature
s=500)
tf_idf_vect.fit(preprocessed_reviews)
```

- **SET 3:** Review text, preprocessed one converted into vectors using (AVG W2v)
- **SET 4:** Review text, preprocessed one converted into vectors using (TFIDF W2v)

### 3. The hyper paramter tuning(find best K)

- Find the best hyper parameter which will give the maximum [AUC](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/receiver-operating-characteristic-curve-roc-curve-and-auc-1/) (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/receiver-operating-characteristic-curve-roc-curve-and-auc-1/>) value
- Find the best hyper paramter using k-fold cross validation or simple cross validation data
- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter tuning

### 4. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure



Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.



Along with plotting ROC curve, you need to print the [confusion matrix](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/) (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/>) with predicted and original labels of test data points



### 5. Conclusion

- You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library [link](http://zetcode.com/python/prettytable/) (<http://zetcode.com/python/prettytable/>).



[[http://]](http://)

### Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method `fit_transform()` on you train data, and apply the method `transform()` on cv/test data.
4. For more details please go through this [link](https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf). (<https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf>)

In [25]:

```
ing KNN brute force

(x_tr,x_cv, y_tr, y_cv):

[]
range(1,30,2):

NeighborsClassifier(n_neighbors=i,algorithm='brute',n_jobs=-1)
cross_val_score(knn,x_tr, y_tr, cv=6, scoring='accuracy',n_jobs=-1)
s.mean()
s.append(d)

2*cv_scores.index(max(cv_scores)) +1
    optimum value of k is :",optimum_K)
NeighborsClassifier(n_neighbors=optimum_K,algorithm='brute',n_jobs=-1)
tr,y_tr)
.predict(x_cv)
uracy of the classifier:",accuracy_score(y_cv,y_pred))
stackoverflow.com/questions/52910061/implementing-roc-curves-for-k-nn-machine-learnin
threshold = roc_curve(y_cv, y_pred)
auc(fpr, tpr)
'Receiver Operating Characteristic')
or, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
(loc = 'lower right')
0, 1], [0, 1], 'r--')
0, 1])
0, 1])
('True Positive Rate')
('False Positive Rate')
'ROC Curve of kNN')

n.metrics import confusion_matrix
on_matrix(y_cv,y_pred)

= ["negative", "positive"]
Dataframe(cm, index = class_label, columns = class_label)
o(df_cm, annot = True, fmt = "d")
'Confusion Matrix')
("Predicted Label")
("True Label")


sification_report(y_cv, y_pred))
```

### [5.1.1] Applying KNN brute force on BOW, SET 1


In [26]:

```
# Please write all the code with proper documentation
x_tr,x_cv, y_tr, y_cv= train_test_split(preprocessed_reviews,final['Score'],test_si
vect = CountVectorizer()
vect= vect.fit(preprocessed_reviews)
x_tr=vect.transform(x_tr)
x_cv= vect.transform(x_cv)
brute_Knn(x_tr,x_cv, y_tr, y_cv)
```

```
1
3
5
7
9
11
13
15
17
19
21
23
25
27
29
the optimum value of k is : 29
accuracy of the classifier: 0.8355497506530515
```



```
[[ 110 1280]
 [ 105 6927]]
```



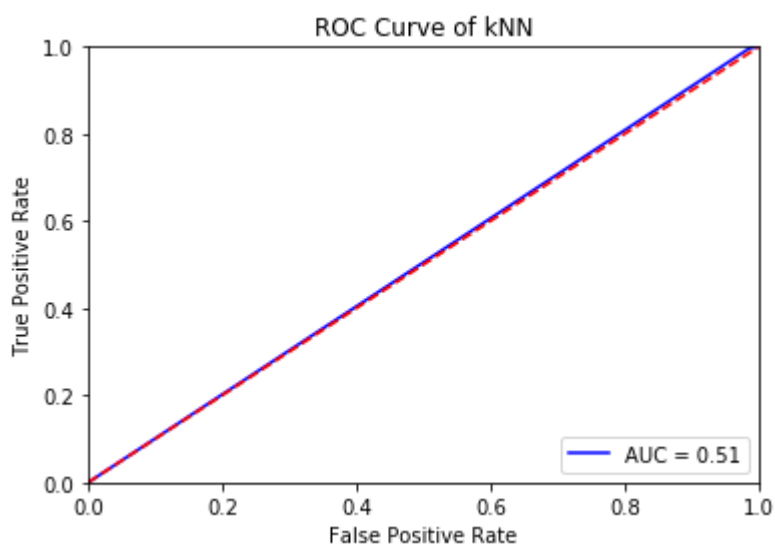
	precision	recall	f1-score	support
0	0.51	0.08	0.14	1390
1	0.84	0.99	0.91	7032
micro avg	0.84	0.84	0.84	8422
macro avg	0.68	0.53	0.52	8422
weighted avg	0.79	0.84	0.78	8422

### [5.1.2] Applying KNN brute force on TFIDF, SET 2


In [27]:

```
# Please write all the code with proper documentation
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tf_idf_vect.fit(preprocessed_reviews)
final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
x_tr,x_cv, y_tr, y_cv= train_test_split(final_tf_idf,final['Score'],test_size=0.3)
brute_Knn(x_tr,x_cv, y_tr, y_cv)
```

```
1
3
5
7
9
11
13
15
17
19
21
23
25
27
29
the optimum value of k is : 1
accuracy of the classifier: 0.8399430064117787
```



```
[[ 18 1327]
 [ 21 7056]]
```



	precision	recall	f1-score	support
0	0.46	0.01	0.03	1345
1	0.84	1.00	0.91	7077
micro avg	0.84	0.84	0.84	8422
macro avg	0.65	0.51	0.47	8422
weighted avg	0.78	0.84	0.77	8422


[5.1.3] Applying KNN brute force on AVG W2V, **SET 3**

In [28]:


```
all the code with proper documentation
nce=[]
in preprocessed_reviews:
    sentence.append(sentence.split())
    w2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    list(w2v_model.wv.vocab)
    = []; # the avg-w2v for each sentence/review is stored in this list
    qdm(list_of_sentence): # for each review/sentence
    = np.zeros(50) # as word vectors are of zero length 50, you might need to change thi
    =0; # num of words with a valid vector in the sentence/review
in sent: # for each word in a review/sentence
    and in w2v_words:
    vec = w2v_model.wv[word]
    sent_vec += vec
    cnt_words += 1
    words != 0:
    vec /= cnt_words
    ors.append(sent_vec)
    tr, y_cv= train_test_split(sent_vectors,final['Score'],test_size=0.3)
    r,x_cv, y_tr, y_cv)
```

100%|██████████| 28072/28072 [01:03<00:00, 439.64it/s]

```
1
3
5
7
9
11
13
15
17
19
21
23
25
27
29
the optimum value of k is : 15
accuracy of the classifier: 0.863571598195203
```



```
[[ 371  999]
 [ 150 6902]]
```



	precision	recall	f1-score	support
0	0.71	0.27	0.39	1370
1	0.87	0.98	0.92	7052
micro avg	0.86	0.86	0.86	8422
macro avg	0.79	0.62	0.66	8422
weighted avg	0.85	0.86	0.84	8422

[5.1.4] Applying KNN brute force on TFIDF W2V, SET 4



In [29]:


```
model = TfidfVectorizer()
tf_idf_matrix = model.fit_transform(preprocessed_reviews)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1
x_tr,x_cv, y_tr, y_cv= train_test_split(tfidf_sent_vectors,final['Score'],test_size=0.2)
brute_Knn(x_tr,x_cv, y_tr, y_cv)
```


100%|██████████| 28072/28072 [10:42<00:00, 43.72it/s]

1  
3  
5  
7  
9  
11  
13  
15  
17  
19  
21  
23  
25  
27  
29

the optimum value of k is : 15  
accuracy of the classifier: 0.8554975065305154



```
[[ 297 1084]
 [ 133 6908]]
```



	precision	recall	f1-score	support
0	0.69	0.22	0.33	1381
1	0.86	0.98	0.92	7041
micro avg	0.86	0.86	0.86	8422
macro avg	0.78	0.60	0.62	8422
weighted avg	0.84	0.86	0.82	8422

In [ ]:

In [30]:

```
## [5.2] Applying KNN kd-tree
def kdtree_Knn(x_tr,x_cv, y_tr, y_cv):

    cv_scores=[]
    for i in range(1,30,2):

        knn= KNeighborsClassifier(n_neighbors=i,algorithm='kd_tree',n_jobs=-1)
        scores=cross_val_score(knn,x_tr, y_tr, cv=6, scoring='accuracy',n_jobs=-1)
        d= round(scores.mean(),2)
        cv_scores.append(d)
        print(i)

    optimum_K=2*cv_scores.index(max(cv_scores)) +1
    print("the optimum value of K is:",optimum_K)
    knn= KNeighborsClassifier(n_neighbors=optimum_K,algorithm='kd_tree',n_jobs=-1)
    knn.fit(x_tr,y_tr)
    y_pred=knn.predict(x_cv)
    print("accuracy of the classifier:",accuracy_score(y_cv,y_pred))
# https://stackoverflow.com/questions/52910061/implementing-roc-curves-for-k-nn-
fpr, tpr, threshold = roc_curve(y_cv, y_pred)
roc_auc = auc(fpr, tpr)
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.title('ROC Curve of kNN')
plt.show()
from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_cv,y_pred)
print(cm)
class_label = ["negative", "positive"]
df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
sns.heatmap(df_cm, annot = True, fmt = "d")
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
from sklearn.metrics import confusion_matrix
print(classification_report(y_cv, y_pred))
```


### [5.2.1] Applying KNN kd-tree on BOW, SET 5

In [31]:


```
# Please write all the code with proper documentation
x_tr,x_cv, y_tr, y_cv= train_test_split(preprocessed_reviews,final['Score'],test_si
vect = CountVectorizer(min_df=10, max_features=500)
x_tr = vect.fit_transform(x_tr)
x_cv= vect.transform(x_cv)
from scipy.sparse import csr_matrix, issparse
x_tr=x_tr.todense()
x_cv=x_cv.todense()

kdtree_Knn(x_tr,x_cv, y_tr, y_cv)
```

```
1
3
5
7
9
11
13
15
17
19
21
23
25
27
29
the optimum value of K is: 9
accuracy of the classifier: 0.8335312277368796
```



```
[[ 334  980]
 [ 422 6686]]
```



	precision	recall	f1-score	support
0	0.44	0.25	0.32	1314
1	0.87	0.94	0.91	7108
micro avg	0.83	0.83	0.83	8422
macro avg	0.66	0.60	0.61	8422
weighted avg	0.81	0.83	0.81	8422

## [5.2.2] Applying KNN kd-tree on TFIDF, SET 6


In [35]:

```
# Please write all the code with proper documentation
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2),min_df=10, max_features=500 )
tf_idf_vect.fit(preprocessed_reviews)
final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
x_tr,x_cv, y_tr, y_cv= train_test_split(final_tf_idf,final['Score'],test_size=0.3)
x_tr=x_tr.todense()
x_cv=x_cv.todense()
kdtree_Knn(x_tr,x_cv, y_tr, y_cv)
```


1  
3  
5  
7  
9  
11  
13  
15  
17  
19  
21  
23  
25  
27  
29

the optimum value of K is: 5

accuracy of the classifier: 0.8420802659700783



```
[[ 84 1250]
 [ 80 7008]]
```



	precision	recall	f1-score	support
0	0.51	0.06	0.11	1334
1	0.85	0.99	0.91	7088
micro avg	0.84	0.84	0.84	8422
macro avg	0.68	0.53	0.51	8422
weighted avg	0.80	0.84	0.79	8422

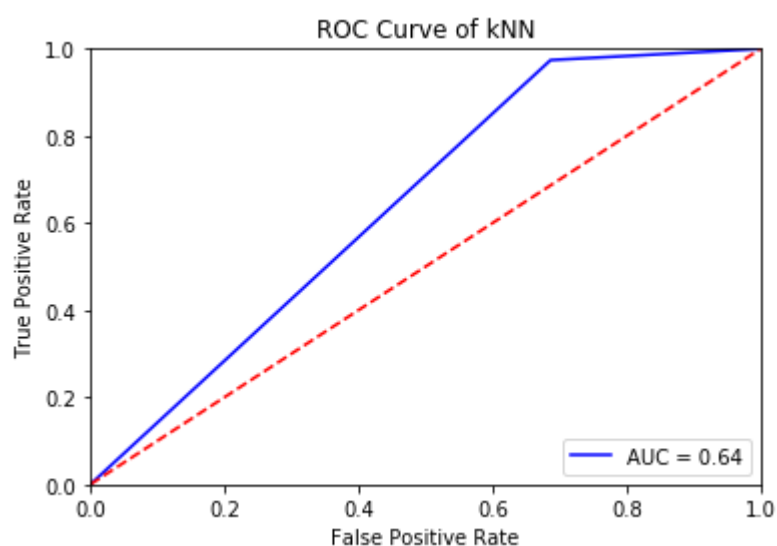
### [5.2.3] Applying KNN kd-tree on AVG W2V, SET 3

In [33]:


```
# Please write all the code with proper documentation
```

```
x_tr,x_cv, y_tr, y_cv= train_test_split(sent_vectors,final['Score'],test_size=0.3)
kdtree_Knn(x_tr,x_cv, y_tr, y_cv)
```

```
1
3
5
7
9
11
13
15
17
19
21
23
25
27
29
the optimum value of K is: 11
accuracy of the classifier: 0.8679648539539302
```



```
[[ 424  929]
 [ 183 6886]]
```






	precision	recall	f1-score	support
0	0.70	0.31	0.43	1353
1	0.88	0.97	0.93	7069
micro avg	0.87	0.87	0.87	8422
macro avg	0.79	0.64	0.68	8422
weighted avg	0.85	0.87	0.85	8422

#### [5.2.4] Applying KNN kd-tree on TFIDF W2V, SET 4


In [34]:

```
# Please write all the code with proper documentation
x_tr,x_cv, y_tr, y_cv= train_test_split(tfidf_sent_vectors,final['Score'],test_size=
kdtree_Knn(x_tr,x_cv, y_tr, y_cv)
```

```
1
3
5
7
9
11
13
15
17
19
21
23
25
27
29
the optimum value of K is: 7
accuracy of the classifier: 0.8576347660888151
```



```
[[ 377  971]
 [ 228 6846]]
```



	precision	recall	f1-score	support
0	0.62	0.28	0.39	1348
1	0.88	0.97	0.92	7074
micro avg	0.86	0.86	0.86	8422
macro avg	0.75	0.62	0.65	8422
weighted avg	0.84	0.86	0.83	8422

## [6] Conclusions

In [8]:

```
# Please compare all your models using Prettytable library
#https://stackoverflow.com/questions/36423259/how-to-use-pretty-table-in-python-to-
from prettytable import PrettyTable
x = PrettyTable(["Vectorizer", "Model", "Hyper parameter", "AUC", "ACCURACY"])
x.add_row(["BOW", "Brute", 29, 0.53, 83])
x.add_row(["TFIDF", "Brute", 1, 0.51, 83])
x.add_row(["AVG W2V", "Brute", 15, 0.62, 86])
x.add_row(["TFIDF W2V", "Brute", 15, 0.60, 85])
x.add_row(["BOW", "KDtree", 9, 0.60, 83])
x.add_row(["TFIDF", "KDtree", 5, 0.53, 84])
x.add_row(["AVG W2V", "KDtree", 11, 0.64, 86])
x.add_row(["TFIDF W2V", "KDtree", 7, 0.62, 85])
print(x)
```

Vectorizer	Model	Hyper parameter	AUC	ACCURACY
BOW	Brute	29	0.53	83
TFIDF	Brute	1	0.51	83
AVG W2V	Brute	15	0.62	86
TFIDF W2V	Brute	15	0.6	85
BOW	KDtree	9	0.6	83
TFIDF	KDtree	5	0.53	84
AVG W2V	KDtree	11	0.64	86
TFIDF W2V	KDtree	7	0.62	85

1. After running all the vectorizers on brute and KD tree models they have accuracy around 85% on test data.
2. If we look into confusion matrix and classification report the classifier performs poor in predicting negative reviews.
3. Even though the classifier has 85% it is not able to predict negative reviews this seems to be overfitting problem
4. The classifier might work well if we take balanced data.
5. AVG W2V vectorizer seems to work well for both KDtree and brute models compared to accuracy of other vectorizers.